




Datanet 2015

Assignment introduction
Kenneth Skovhede


+ Overview



- 4 assignments
- Each assignment scores 0-10 points
- You need 24 points to qualify for the exam
- No re-submissions


+ Structure of assignments

1. Read/research
2. Implement
3. Write report
4. Hand in




+ 1. Read and research

- Read the assignment
- Find required extra material
- Make a hypothesis for the assignment
 - I.e. describe what you are examining
- Write it down






2. Implementation



- Make an implementation that allows you test your hypothesis
- Most likely you will have to go back and research
- You do not get points for source code!
 - But try to make it readable so we can verify your claims
- The implementation is a tool to help you understand the details, complexity and imprecisions of a real-world example




3. Report



- You get points for the report!
 - Read the “your report must contain ...” section!
- Your report must be in the ACM standard template
- Try to use this format:
 - Abstract
 - Introduction
 - Design
 - Implementation
 - Experiments
 - Results



4. Handing in



■ Deadlines


- 1st assignment: 2015-04-28 23:55 (11:55pm)
- 2nd assignment: 2015-05-08 23:55 (11:55pm)
- 3rd assignment: 2015-05-26 23:55 (11:55pm)
- 4th assignment: 2015-06-04 23:55 (11:55pm)

■ Plagiarism

- I check if you copy stuff (i.e. from Wikipedia)
- If you copy something, make sure you quote (or reference)
- You may work in groups, but be sure you hand in your own report




1st assignment: HTTP Client



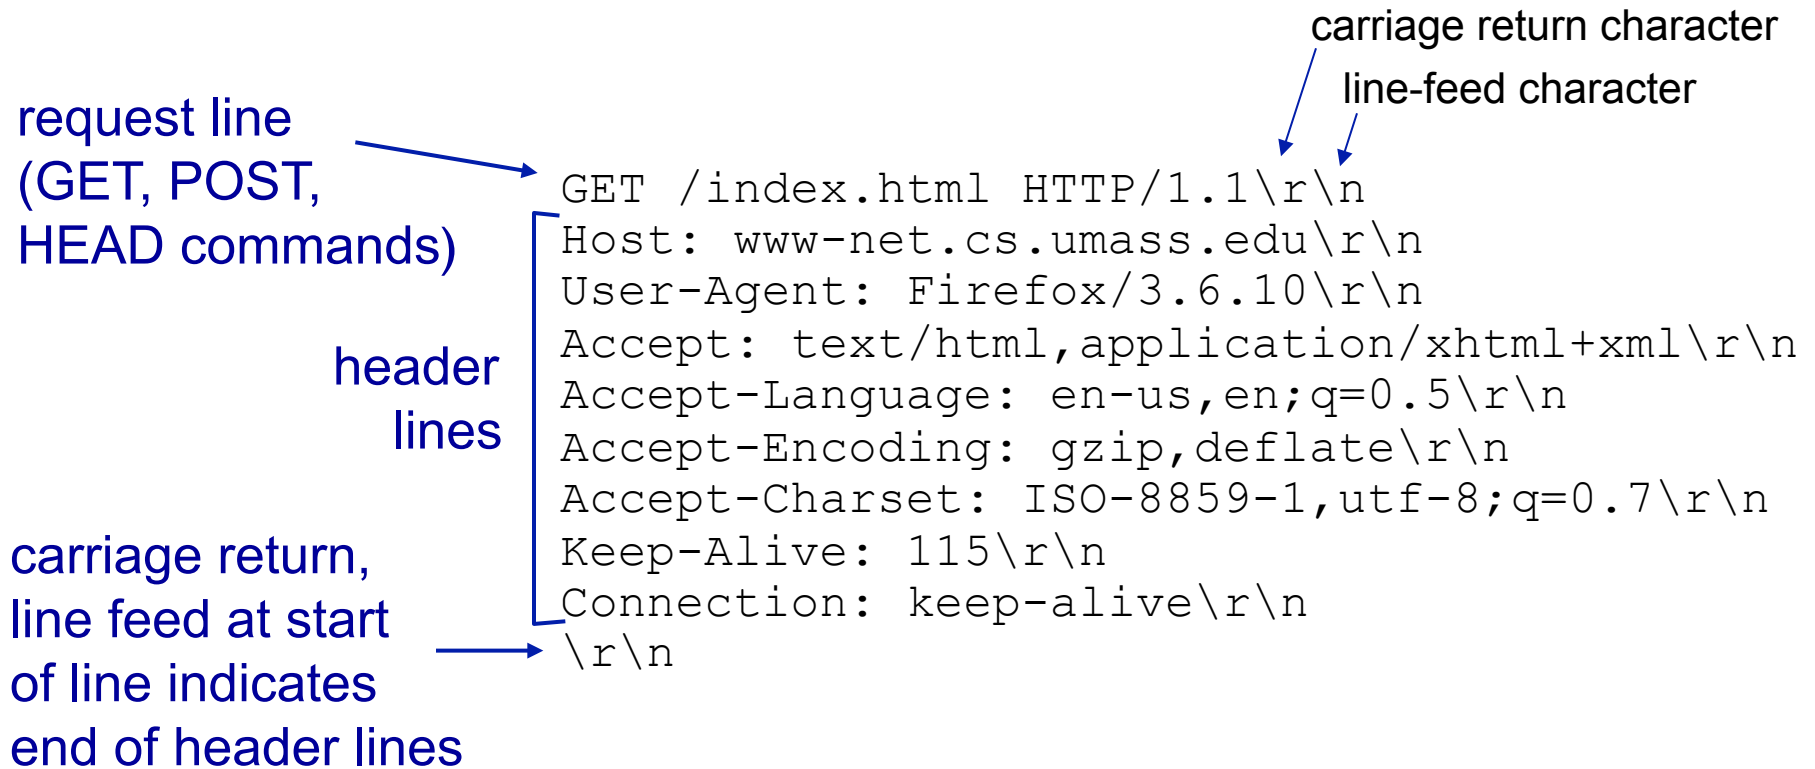
- Build a rough version of `wget` / `curl`
 - Learn sockets in practice
 - Learn HTTP in practice
 - Measure latency, topology and bandwidth

+ URLs



- Many features in a string:
 - `scheme://user:password@host:port/path?query_string#fragment_id`
- Most are optional:
 - `scheme://host/path`
- An example:
 - `http://www.diku.dk/index.html`

+ HTTP request message

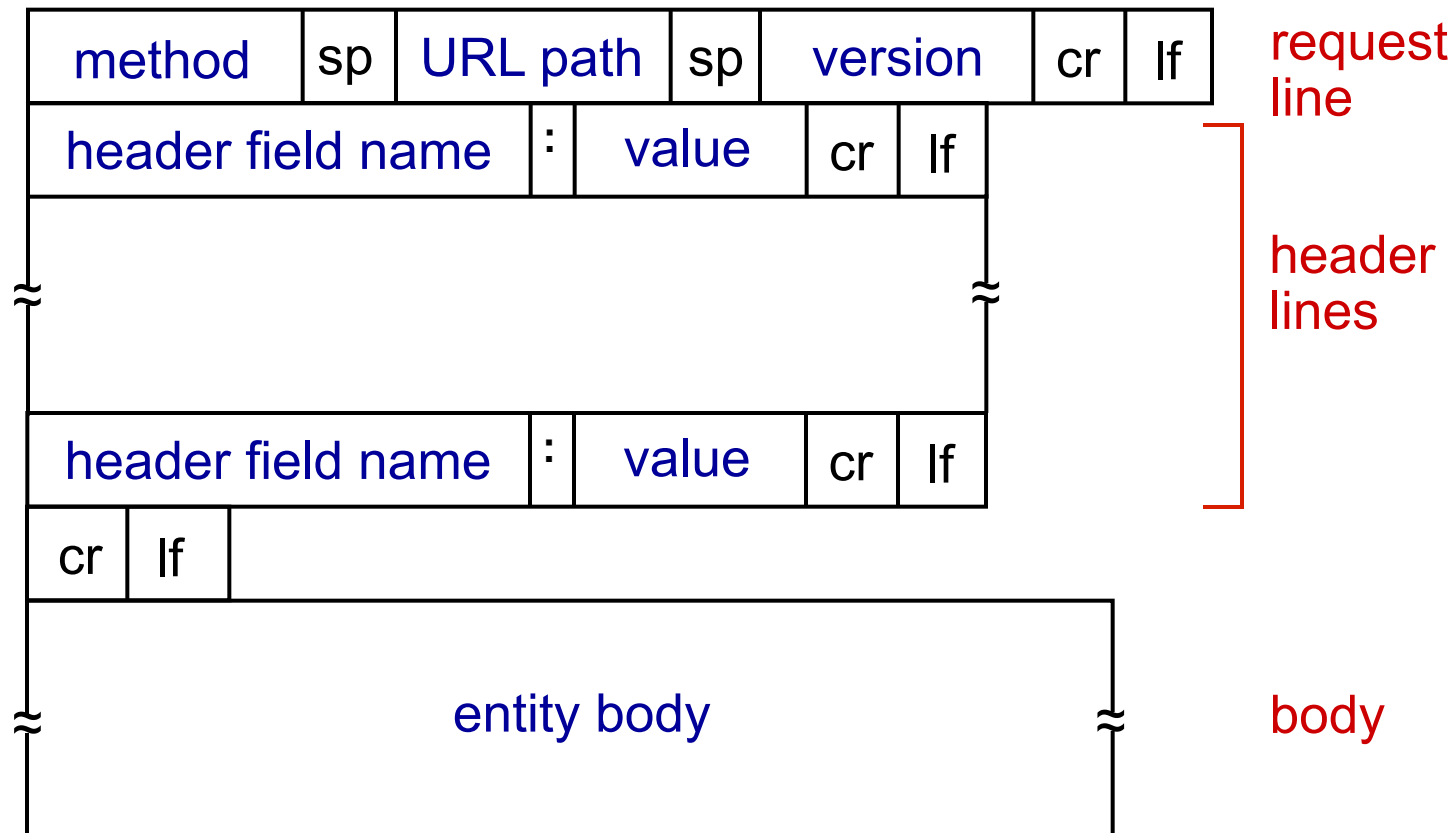


The diagram illustrates the structure of an HTTP request message. It consists of a request line followed by header lines, and a final blank line. Annotations with arrows point to specific parts of the message:

- request line (GET, POST, HEAD commands)**: Points to the first line of the message: `GET /index.html HTTP/1.1\r\n`.
- header lines**: Points to the block of lines between the request line and the final blank line, including `Host: www-net.cs.umass.edu\r\n`, `User-Agent: Firefox/3.6.10\r\n`, `Accept: text/html,application/xhtml+xml\r\n`, `Accept-Language: en-us,en;q=0.5\r\n`, `Accept-Encoding: gzip,deflate\r\n`, `Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n`, `Keep-Alive: 115\r\n`, and `Connection: keep-alive\r\n`.
- carriage return character**: Points to the `\r` character in the first line.
- line-feed character**: Points to the `\n` character in the first line.
- carriage return, line feed at start of line indicates end of header lines**: Points to the final blank line `\r\n`.

```
GET /index.html HTTP/1.1\r\nHost: www-net.cs.umass.edu\r\nUser-Agent: Firefox/3.6.10\r\nAccept: text/html,application/xhtml+xml\r\nAccept-Language: en-us,en;q=0.5\r\nAccept-Encoding: gzip,deflate\r\nAccept-Charset: ISO-8859-1,utf-8;q=0.7\r\nKeep-Alive: 115\r\nConnection: keep-alive\r\n\r\n
```

+ HTTP request message: general format



+ Sockets

■ An example in Python:

```
import socket
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.connect(("www.diku.dk", 80))
s.send(
    "GET / HTTP/1.1\r\n" +
    "Host: www.diku.dk\r\n" +
    "Connection: close\r\n" +
    "\r\n")

msg = s.recv(4096)
while msg != '':
    print msg
    msg = s.recv(4096)
s.close()
```

+ HTTP Response Message

status line
(protocol
status code
status phrase)

header
lines


data, e.g.,
requested
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02 GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html;\r\n
    charset=ISO-8859-1\r\n
\r\n
data data data data data ...
```




Too easy?

+ Use the time wisely



- Time to learn LaTeX?
- Familiarize yourself with the ACM template
- Think in academic writing style
 - Introduction
 - Hypothesis
 - Experiments
 - Results
 - Conclusion
- Try out Amazon EC2

+ Collaboration rules



- You may work in groups of 2-4 persons
- But you must hand in an *individual* report
- If you cite/copy/reference something, you *MUST* state the source, i.e.:
 - "Goats are among the earliest animals domesticated by humans." [Wikipedia]
- Failure to follow this is considered cheating, and may have repercussions



Questions?